

FinalProject

December 10, 2019

```
[218]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import calendar
from datetime import datetime
from pandas import Series
from math import ceil
import warnings
warnings.filterwarnings('ignore')
df = pd.read_csv('Video_Games_Sales.csv')
```

```
[218]: Name          object
Platform         object
Year_of_Release  float64
Genre            object
Publisher        object
NA_Sales         float64
EU_Sales         float64
JP_Sales         float64
Other_Sales      float64
Global_Sales     float64
Critic_Score     float64
Critic_Count     float64
User_Score       object
User_Count       float64
Developer        object
Rating           object
dtype: object
```

0.1 Data Analysis

Missing Values Show the missing values distribution:

```
[217]: plt.figure(figsize=(14,7))
sns.heatmap(df.isnull(),yticklabels=False,cbar=False,cmap='viridis')
df['Global_Sales_Log'] = np.log1p(df['Global_Sales'])
```

```

df['NA_Sales_Log'] = np.log1p(df['NA_Sales'])
df['EU_Sales_Log'] = np.log1p(df['EU_Sales'])
df['JP_Sales_Log'] = np.log1p(df['JP_Sales'])
df['Other_Sales_Log'] = np.log1p(df['Other_Sales'])

null_columns=df.columns[df.isnull().any()]
df[null_columns].isnull().sum()


```

```

[217]: Year_of_Release      float64
       NA_Sales            float64
       EU_Sales            float64
       JP_Sales            float64
       Other_Sales         float64
       Global_Sales        float64
       Platform_XB          float64
       Platform_X360        float64
       Platform_XOne        float64
       Platform_PC          float64
       Platform_PS          float64
       Platform_PS2         float64
       Platform_PS3         float64
       Platform_PS4         float64
       Platform_PSP         float64
       Platform_PSV         float64
       Platform_GB          float64
       Platform_GBA         float64
       Platform_DS          float64
       Platform_3DS         float64
       Platform_NES         float64
       Platform_SNES        float64
       Platform_N64         float64
       Platform_GC          float64
       Platform_Wii         float64
       Platform_WiiU        float64
       Genre_Action         float64
       Genre_Adventure      float64
       Genre_Fighting       float64
       Genre_Misc           float64
       ...
       Flying Lab Software   float64
       1C, Ino-Co, 1C Company float64
       Autumn Moon          float64
       Tate Interactive      float64
       Elixir Studios        float64
       Compulsion Games      float64
       Camouflaj, LLC        float64
       DMA Design, Rockstar North float64

```

King of the Jungle	float64
Spidersoft, Spiders	float64
Tecmo, Graphic Research	float64
Battlefront.com, 1C, 1C Company	float64
EA Phenomic	float64
Empty Clip Studios	float64
Headgate	float64
Coffee Stain Studios	float64
Boston Animation	float64
React Games	float64
Inferno Games	float64
Katauri Interactive	float64
High Moon Studios, Mercenary Technologies	float64
Infinite Dreams, Paragon 5	float64
Big Red Software	float64
Fluid Studios	float64
Atomic Games	float64
Global_Sales_Log	float64
NA_Sales_Log	float64
EU_Sales_Log	float64
JP_Sales_Log	float64
Other_Sales_Log	float64
Length: 1491, dtype: object	



```
[186]: # we can see there are lots of missing values in Critic_Score, Critic_Count,
#User_Score etc. While there are ONLY TWO missing in Name and Genre.
```

1 Data Cleaning

```
[187]: df.drop(index=[659,14246],inplace=True)
df[df.Year_of_Release.isnull()]
df.Year_of_Release = df.Year_of_Release.fillna(0)
df[df['Year_of_Release']==0].Platform.unique()

PS2_median = df[df['Platform'] == 'PS2']['Year_of_Release'].median()
Wii_median = df[df['Platform'] == 'Wii']['Year_of_Release'].median()
x2600median = df[df['Platform'] == '2600']['Year_of_Release'].median()
X360_median = df[df['Platform'] == 'X360']['Year_of_Release'].median()
GBA_median = df[df['Platform'] == 'GBA']['Year_of_Release'].median()
PC_median = df[df['Platform'] == 'PC']['Year_of_Release'].median()
PS3_median = df[df['Platform'] == 'PS3']['Year_of_Release'].median()
PS_median = df[df['Platform'] == 'PS']['Year_of_Release'].median()
PSP_median = df[df['Platform'] == 'PSP']['Year_of_Release'].median()
XB_median = df[df['Platform'] == 'XB']['Year_of_Release'].median()
GB_median = df[df['Platform'] == 'GB']['Year_of_Release'].median()
DS_median = df[df['Platform'] == 'DS']['Year_of_Release'].median()
x3DS_median = df[df['Platform'] == '3DS']['Year_of_Release'].median()
N64_median = df[df['Platform'] == 'N64']['Year_of_Release'].median()
PSV_median = df[df['Platform'] == 'PSV']['Year_of_Release'].median()
GC_median = df[df['Platform'] == 'GC']['Year_of_Release'].median()

# Function that returns the median of the platform if year = 0. Else it returns
→ the year.
def year_filler(x):

    if x.Year_of_Release == 0:
        if x.Platform == 'PS2':
            return PS2_median
        elif x.Platform == 'Wii':
            return Wii_median
        elif x.Platform == '2600':
            return x2600median
        elif x.Platform == 'X360':
            return X360_median
        elif x.Platform == 'GBA':
            return GBA_median
        elif x.Platform == 'PC':
            return PC_median
```

```

        elif x.Platform == 'PS3':
            return PS3_median
        elif x.Platform == 'PS':
            return PS_median
        elif x.Platform == 'PSP':
            return PSP_median
        elif x.Platform == 'XB':
            return XB_median
        elif x.Platform == 'GB':
            return GB_median
        elif x.Platform == 'DS':
            return DS_median
        elif x.Platform == '3DS':
            return x3DS_median
        elif x.Platform == 'N64':
            return N64_median
        elif x.Platform == 'PSV':
            return PSV_median
        elif x.Platform == 'GC':
            return GC_median
        else:
            return 1900
    else:
        return x.Year_of_Release

# apply function to replace values
df.Year_of_Release = df.apply(year_filler, axis=1)

df['Publisher'].fillna(value='Unknown', inplace=True)
# We will start by filling the missing values of Critic_Score and Critic_Count.
df.Critic_Score = df.Critic_Score.fillna(df.Critic_Score.median())
df.Critic_Count = df.Critic_Count.fillna(df.Critic_Count.median())

# We will start by filling the missing values of Critic_Score and Critic_Count.
# Because Critic_Score has NaN values, we will first fill those NaN values with
→ 0.

df.User_Score = df.User_Score.fillna(0)

# Some values of Critic_Score have the value 'tbd'. These values need to be
→ replaced with 100.
df.User_Score.replace(to_replace='tbd', value=100, inplace=True)

# We will do the same for User_Count
df.User_Count = df.User_Count.fillna(df.User_Count.median())

```

```

# Replace the missing values of Rating with Unknown.
df.Rating.fillna('Unknown',inplace=True)

# Because Developer has too many unique values, we will drop this column.
df.drop('Developer',axis=1,inplace=True)

columns = df.columns
percent_missing = df.isnull().sum() * 100 / len(df)
missing_value_df = pd.DataFrame({'column_name': columns,
                                'percent_missing': percent_missing})

print(missing_value_df)

### Missing Values Show the missing values distribution:
plt.figure(figsize=(14,7))
sns.heatmap(df.isnull(),yticklabels=False,cbar=False,cmap='viridis')
df['Global_Sales_Log'] = np.log1p(df['Global_Sales'])

df['NA_Sales_Log'] = np.log1p(df['NA_Sales'])
df['EU_Sales_Log'] = np.log1p(df['EU_Sales'])
df['JP_Sales_Log'] = np.log1p(df['JP_Sales'])
df['Other_Sales_Log'] = np.log1p(df['Other_Sales'])


null_columns=df.columns[df.isnull().any()]
df[null_columns].isnull().sum()

dforiginal=df;

```

	column_name	percent_missing
Name	Name	0.0
Platform	Platform	0.0
Year_of_Release	Year_of_Release	0.0
Genre	Genre	0.0
Publisher	Publisher	0.0
NA_Sales	NA_Sales	0.0
EU_Sales	EU_Sales	0.0
JP_Sales	JP_Sales	0.0
Other_Sales	Other_Sales	0.0
Global_Sales	Global_Sales	0.0
Critic_Score	Critic_Score	0.0
Critic_Count	Critic_Count	0.0
User_Score	User_Score	0.0
User_Count	User_Count	0.0
Rating	Rating	0.0
Global_Sales_Log	Global_Sales_Log	0.0
NA_Sales_Log	NA_Sales_Log	0.0

EU_Sales_Log	EU_Sales_Log	0.0
JP_Sales_Log	JP_Sales_Log	0.0
Other_Sales_Log	Other_Sales_Log	0.0



2 Data Analysis & Data visualization:

```
[188]: #Plot Top values in the dataset By platform, developer and genre.
import matplotlib.pyplot as plt
%matplotlib inline

# set font
plt.rcParams['font.family'] = 'sans-serif'
plt.rcParams['font.sans-serif'] = 'Helvetica'

# set the style of the axes and the text color
plt.rcParams['axes.edgecolor']='#333F4B'
plt.rcParams['axes.linewidth']=0.8
plt.rcParams['xtick.color']='#b5651d'
plt.rcParams['ytick.color']='#b5651d'
plt.rcParams['text.color']='#b5651d'
plt.rcParams['figure.figsize'] = 15, 5
```

```

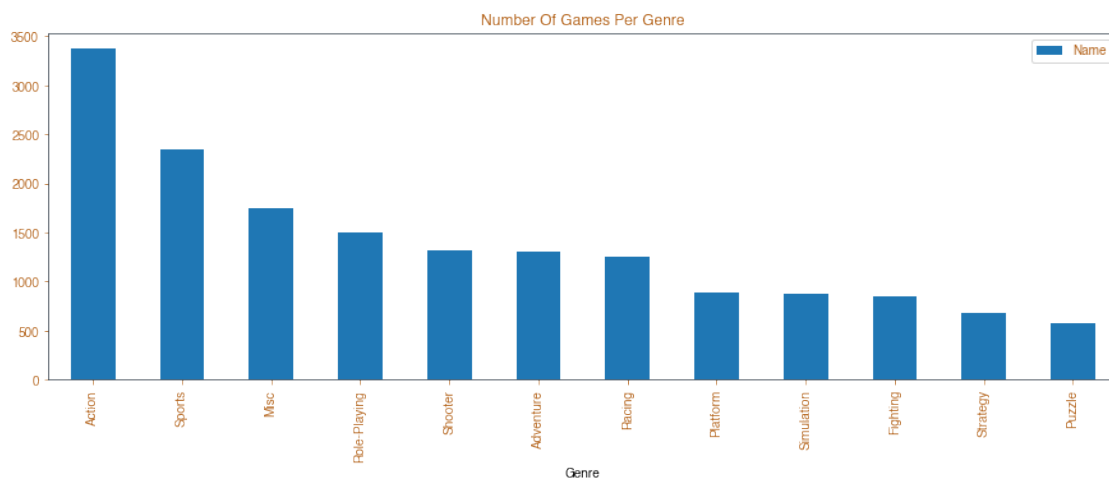

namePlatformDf = df[['Name', 'Platform']].groupby(['Platform']).count().
    ↳sort_values('Name', ascending=False).reset_index()
namePlatformDf.plot(kind='bar', x="Platform", y="Name",title="Number Of Games_
    ↳For Different Console")

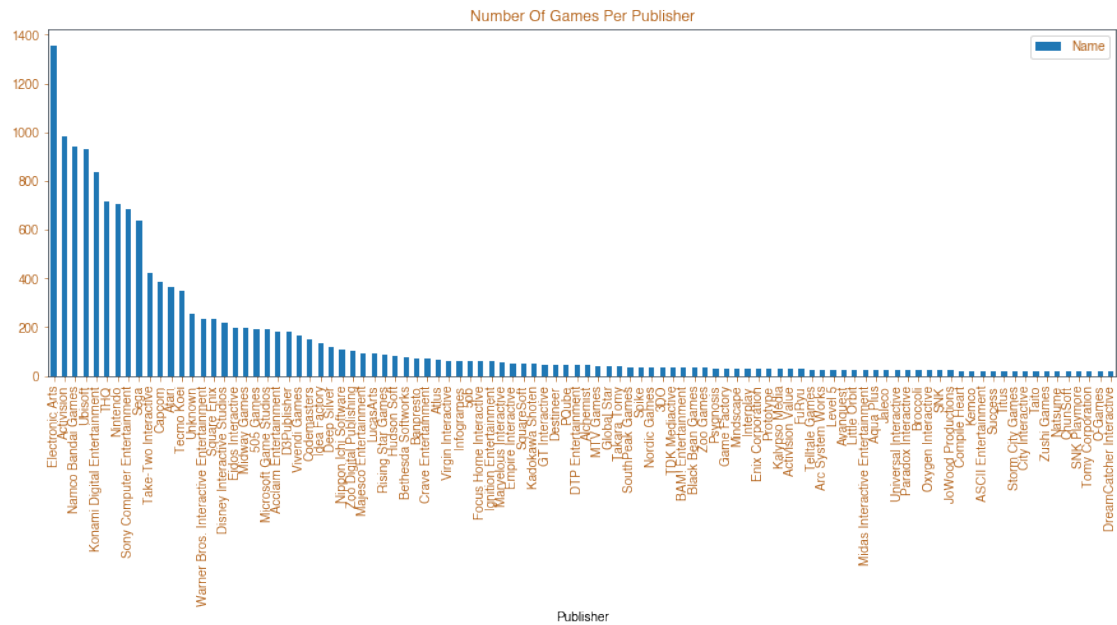
nameGenreDf = df[['Name', 'Genre']].groupby(['Genre']).count().
    ↳sort_values('Name', ascending=False).reset_index()
nameGenreDf.plot(kind='bar', x="Genre", y="Name",title="Number Of Games Per_
    ↳Genre")

nameDeveloperDf = df[['Name', 'Publisher']].groupby(['Publisher']).count().
    ↳sort_values('Name', ascending=False).reset_index()
nameDeveloperDf[:100].plot(kind='bar', x="Publisher", y="Name",title="Number Of_
    ↳Games Per Publisher")

```

[188]: <matplotlib.axes._subplots.AxesSubplot at 0x1a23eeacc0>






DATA CORRELATION

```
[189]: plt.figure(figsize=(12, 8))
```

```
vg_corr = df.corr()
sns.heatmap(vg_corr,
             xticklabels = vg_corr.columns.values,
             yticklabels = vg_corr.columns.values,
             annot = True);
```



```
[190]: ## UTILITY FUNCTION FOR PRINTING CLASSIFICATION REPORT
```

```
[191]: import seaborn as sns
import numpy as np
from sklearn.metrics import precision_recall_fscore_support
import matplotlib.pyplot as plt

y = np.random.randint(low=0, high=10, size=100)
y_p = np.random.randint(low=0, high=10, size=100)

def plot_classification_report(y_tru, y_prd, figsize=(10, 10), ax=None):

    plt.figure(figsize=figsize)

    xticks = ['precision', 'recall', 'f1-score', 'support']
    yticks = list(np.unique(y_tru))
    yticks += ['avg']

    rep = np.array(precision_recall_fscore_support(y_tru, y_prd)).T
    avg = np.mean(rep, axis=0)
    avg[-1] = np.sum(rep[:, -1])
    rep = np.insert(rep, rep.shape[0], avg, axis=0)
```

```
sns.heatmap(rep,
             annot=True,
             cbar=False,
             xticklabels=xticks,
             yticklabels=yticks,
             ax=ax)
```

[192]: *## UTILITY FUNCTION FOR PRINTING CONFUSION MATRIX*

```
[193]: import itertools
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import confusion_matrix

# Source: http://scikit-learn.org/stable/auto\_examples/model\_selection/
#         plot\_confusion\_matrix.html#confusion-matrix

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")
    plt.tight_layout()
```

```
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

[194]: *# Data Transformation*

```
df.drop(index=[0], inplace=True)

# Transform datatypes change the datatype from real number -> integer
# Transform year from float to integer
df['Year_of_Release'] = df.Year_of_Release.astype(int)

# Transform Critic_Score from float to integer
df['Critic_Score'] = df.Critic_Score.astype(int)

# Transform Critic_Count from float to integer
df['Critic_Count'] = df.Critic_Count.astype(int)

# Transform User_Count from float to integer
df['User_Count'] = df.User_Count.astype(int)

# Transform User_Score to int
df.User_Score = pd.to_numeric(df.User_Score, errors='coerce')
```

[195]: *# Categorical encoding and standardization*
First we will create a new feature, based on the year of release.
The new feature 'years_after_release' = year of release - release date of the
→ platform

```
# Create new, empty column
df['year_after_release'] = ''

# Create variables for all console release years
PS2_release = 2000
Wii_release = 2006
x2600_release = 1977
X360_release = 2005
GBA_release = 2001
PS3_release = 2006
PS_release = 1994
PSP_release = 2004
XB_release = 2001
GB_release = 1989
DS_release = 2004
x3DS_release = 2011
N64_release = 1996
PSV_release = 2011
GC_release = 2001
```

```
def year_after_release_filler(x):

    if x.Platform == 'PS2':
        return x.Year_of_Release - PS2_release
    elif x.Platform == 'Wii':
        return x.Year_of_Release - Wii_release
    elif x.Platform == '2600':
        return x.Year_of_Release - x2600_release
    elif x.Platform == 'X360':
        return x.Year_of_Release - X360_release
    elif x.Platform == 'GBA':
        return x.Year_of_Release - GBA_release
    elif x.Platform == 'PS3':
        return x.Year_of_Release - PS3_release
    elif x.Platform == 'PS':
        return x.Year_of_Release - PS_release
    elif x.Platform == 'PSP':
        return x.Year_of_Release - PSP_release
    elif x.Platform == 'XB':
        return x.Year_of_Release - XB_release
    elif x.Platform == 'GB':
        return x.Year_of_Release - GB_release
    elif x.Platform == 'DS':
        return x.Year_of_Release - DS_release
    elif x.Platform == '3DS':
        return x.Year_of_Release - x3DS_release
    elif x.Platform == 'N64':
        return x.Year_of_Release - N64_release
    elif x.Platform == 'PSV':
        return x.Year_of_Release - PSV_release
    elif x.Platform == 'GC':
        return x.Year_of_Release - GC_release
    else:
        return 1
```

```
df.year_after_release = df.apply(year_after_release_filler, axis=1)
df.drop(index=[15959], inplace=True)
df[df['year_after_release'] < 0]
```

```
[195]:
```

	Name	Platform	Year_of_Release	\
1340	Disney's DuckTales	GB	1988	
2076	NFL Fever 2002	XB	2000	
12301	ESPN Winter X-Games: Snowboarding 2002	GBA	2000	

	Genre		Publisher	NA_Sales	EU_Sales	JP_Sales	\
1340	Platform		Capcom	0.82	0.23	0.35	
2076	Sports		Microsoft Game Studios	0.74	0.21	0.00	
12301	Sports		Konami Digital Entertainment	0.05	0.02	0.00	

	Other_Sales	Global_Sales	...	Critic_Count	User_Score	User_Count	\
1340	0.03	1.43	...	21	0.0	24	
2076	0.04	0.99	...	24	8.5	10	
12301	0.00	0.06	...	21	0.0	24	

	Rating	Global_Sales_Log	NA_Sales_Log	EU_Sales_Log	JP_Sales_Log	\
1340	Unknown	0.887891	0.598837	0.207014	0.300105	
2076	E	0.688135	0.553885	0.190620	0.000000	
12301	Unknown	0.058269	0.048790	0.019803	0.000000	

	Other_Sales_Log	year_after_release
1340	0.029559	-1
2076	0.039221	-1
12301	0.000000	-1

[3 rows x 21 columns]

2.0.1 Learning(Training) and Evaluation(Testing)

2.0.2 Split Train Cases and Test Cases

```
[196]: # Split training-test data
# For the first ML models, we will use 'Global_Sales_Log' as the target
#       variable.

# Standardizing numerical features
# User_Score features that need to be standardized are: Critic_Score,
#       Critic_Count, User_Count, year_after_release

df10=df

# Encoding categorical features
# We will now create dummy variables for Platform, Genre, Publisher and Rating.

df_dummies = pd.get_dummies(df10[['Platform','Genre','Publisher','Rating']],
#       drop_first = True)
# Merge the dummies dataframe with the original dataframe
df10 = pd.merge(df_dummies, df10, left_index=True, right_index=True)

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

```

print(df10.columns)
scaled_features = scaler.fit_transform(df10[['Critic_Score',
→ 'Critic_Count', 'User_Score', 'User_Count', 'year_after_release']])
scaled_df = pd.DataFrame(scaled_features, columns=['Critic_Score',
→ 'Critic_Count', 'User_Score', 'User_Count', 'year_after_release'])

df10 = pd.merge(scaled_df, df, left_index=True, right_index=True)

# # Drop original non-standardized features
columns = ['Critic_Score_y',
→ 'Critic_Count_y', 'User_Score_y', 'User_Count_y', 'NA_Sales',
→ 'EU_Sales', 'JP_Sales', 'Other_Sales', 'Global_Sales', 'year_after_release_y']
df10.drop(columns, inplace=True, axis=1)

X = df10.
→ drop(['Global_Sales_Log', 'NA_Sales_Log', 'EU_Sales_Log', 'JP_Sales_Log', 'Other_Sales_Log'],
→ axis=1)
y = df10['Global_Sales_Log']
X=df10[['Critic_Score_x', 'Critic_Count_x', 'User_Score_x', 'User_Count_x',
→ 'year_after_release_x', 'NA_Sales_Log']]

```

```

Index(['Platform_3D0', 'Platform_3DS', 'Platform_DC', 'Platform_DS',
      'Platform_GB', 'Platform_GBA', 'Platform_GC', 'Platform_GEN',
      'Platform_GG', 'Platform_N64',
      ...,
      'Critic_Count', 'User_Score', 'User_Count', 'Rating',
      'Global_Sales_Log', 'NA_Sales_Log', 'EU_Sales_Log', 'JP_Sales_Log',
      'Other_Sales_Log', 'year_after_release'],
      dtype='object', length=649)

```

[197]: *##Learning and evaluation*

```

from sklearn.model_selection import KFold
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
→ random_state=101)

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import learning_curve

```

2.1 1.First ML Problem

X=> Critic_Score_x', 'Critic_Count_x', 'User_Score_x', 'User_Count_x', 'year_after_release_x'

y=> Global_Sales_Log

1.Linear Regression

```
[222]: # 1.Linear Regression

from sklearn.linear_model import LinearRegression


lr = LinearRegression()
lr.fit(X_train,y_train)
lr_predictions = lr.predict(X_test)
# Evaluate model
from sklearn.metrics import explained_variance_score, mean_squared_error, \
    mean_absolute_error

sns.scatterplot(lr_predictions,y_test)

MAE_lr = mean_absolute_error(y_test, lr_predictions)
MSE_lr = mean_squared_error(y_test, lr_predictions)
var_lr = explained_variance_score(y_test, lr_predictions)

print("accuracy :" + str(lr.score(X_test,y_test)))
print("mean_absolute_error :" + str(MAE_lr))
print("mean_squared_error :" + str(MSE_lr))
print("variance_score :" + str(var_lr))
```

```
accuracy :0.8523184660252421
mean_absolute_error :0.09902835743695935
mean_squared_error :0.02286022063308966
variance_score :0.8523184741545611
```



2. Ridge Regression

```
[223]: # 2. Ridge Regression

from sklearn.linear_model import Ridge

ridge = Ridge()
parameters = {'alpha':[0.001,0.005,0.01,0.1,0.5,1], 'normalize':[True,False],
              →'tol':[1e-06,5e-06,1e-05,5e-05]}
grid_ridge = GridSearchCV(ridge, parameters, cv=10, verbose=1, scoring =
              →'explained_variance')
grid_ridge.fit(X_train, y_train)

print(grid_ridge.best_score_)
print(grid_ridge.best_params_)

ridge_optimized = Ridge(alpha= 1, normalize= False, tol=1e-06)

ridge_optimized.fit(X_train,y_train)

ridge_pred = ridge_optimized.predict(X_test)

MAE_ridge = mean_absolute_error(y_test, ridge_pred)
MSE_ridge = mean_squared_error(y_test, ridge_pred)
var_ridge = explained_variance_score(y_test, ridge_pred)
print("mean_absolute_error :" + str(MAE_ridge))
print("mean_squared_error :" + str(MSE_ridge))
print("variance :" + str(var_ridge))
print("accuracy :" + str(ridge_optimized.score(X_test,y_test)))
sns.scatterplot(ridge_pred,y_test)
```


Fitting 10 folds for each of 48 candidates, totalling 480 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

```
0.8434304695922168
{'alpha': 0.1, 'normalize': False, 'tol': 1e-06}
mean_absolute_error :0.09908128080265304
mean_squared_error :0.022863384030535606
variance :0.8522980350467355
accuracy :0.8522980298538295
```

[Parallel(n_jobs=1)]: Done 480 out of 480 | elapsed: 1.3s finished

[223]: <matplotlib.axes._subplots.AxesSubplot at 0x1a237acb38>




3. RandomForest Regression

```
[226]: from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor()
rf.fit(X_train,y_train)
rf_predictions = rf.predict(X_test)
MAE_rf = mean_absolute_error(y_test, rf_predictions)
MSE_rf = mean_squared_error(y_test, rf_predictions)
print("mae: " + str(MAE_rf))
print("mse: " + str(MSE_rf))
print("accuracy :" + str(rf.score(X_test,y_test)))
print("variance :" + str(explained_variance_score(y_test, rf_predictions)))
```

```
mae: 0.09884309984717227
mse: 0.02510867527810916
accuracy :0.8377930055592785
variance :0.8379337860864112
```

```
[201]: sns.scatterplot(rf_predictions,y_test)
```

```
[201]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1be7a2b0>
```




2.1.1 Evaluate different models

```
[202]: # # Comparison of MSE of the models
MSEs = [0.022860,0.022863,0.02479]

models = ['Linear Regression','Ridge Regression','Random Forest Regression']
plt.figure(figsize=(14,6))
sns.barplot(models,MSEs,palette='husl')
```

```
[202]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1cc63d68>
```



2.1.2 2. Second ML Problem

X=> Year_of_Release Critic_Score

y=> Hit

2.1.3 RandomForestClassifier,LogisticRegression

```
[231]: from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, f1_score, accuracy_score, \
    confusion_matrix
from sklearn import svm
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import classification_report
```

```

dfa= df.copy()
dfb =_
    →dfa[['Name', 'Platform', 'Genre', 'Publisher', 'Year_of_Release', 'Critic_Score', 'Global_Sales']]
dfb = dfb.dropna().reset_index(drop=True)
df2 =_
    →dfb[['Platform', 'Genre', 'Publisher', 'Year_of_Release', 'Critic_Score', 'Global_Sales']]
df2['Hit'] = df2['Global_Sales']
df2.drop('Global_Sales', axis=1, inplace=True)

def hit(sales):
    if sales >= 1:
        return 1
    else:
        return 0

df2['Hit'] = df2['Hit'].apply(lambda x: hit(x))

df3 = df2[['Year_of_Release', 'Critic_Score', 'Hit']]
y = df3['Hit'].values
df3 = df3.drop(['Hit'], axis=1)
X = df3.values

Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, test_size=0.50,_
    →random_state=2)

print('X Year_of_Release Critic_Score')
print('Y Hit')

radm = RandomForestClassifier(random_state=2).fit(Xtrain, ytrain)
y_val_1 = radm.predict_proba(Xtest)
ypred3 = radm.predict(Xtest)

print("Validation accuracy with RandomForestClassifier: %0.2f"%
    →%(accuracy_score(ytest, ypred3)))

MAE_rf = mean_absolute_error(ytest, ypred3)
MSE_rf = mean_squared_error(ytest, ypred3)
print("mae for RandomForestClassifier: " + str(MAE_rf))
print("mse for RandomForestClassifier: " + str(MSE_rf))
print("accuracy for RandomForestClassifier:" + str(radm.score(Xtest,ytest)))
print("variance for RandomForestClassifier : " +_
    →str(explained_variance_score(ytest, ypred3)))

log_reg = LogisticRegression().fit(Xtrain, ytrain)
y_val_2 = log_reg.predict_proba(Xtest)
ypred = log_reg.predict(Xtest)

```

```

print("Validation accuracy with LogisticRegression: %.2f"%
      →%(accuracy_score(ytest, ypred)))

print('Regression on training set:',log_reg.score(Xtrain, ytrain))
print('Regression Score on test set:',log_reg.score(Xtest, ytest))

print("mae for LogisticRegression: " + str(mean_absolute_error(ytest, ypred)))
print("mse for LogisticRegression: " + str(mean_squared_error(ytest, ypred)))
print("accuracy for LogisticRegression: " + str(log_reg.score(Xtest,ytest)))
print("variance for LogisticRegression : " + str(explained_variance_score(ytest,
      →ypred)))

from sklearn.linear_model import LinearRegression

LR = LinearRegression()
model = LR.fit(Xtrain, ytrain)
pd.DataFrame({'features': df3.columns, 'estimatedCoefficients': model.
      →coef_})[['features', 'estimatedCoefficients']].
      →sort_values(by='estimatedCoefficients', ascending=False)

```

```

X Year_of_Release Critic_Score
Y Hit
Validation accuracy with RandomForestClassifier: 0.82
mae for RandomForestClassifier: 0.18015534953645704
mse for RandomForestClassifier: 0.18015534953645704
accuracy for RandomForestClassifier:0.8198446504635429
variance for RandomForestClassifier :-0.20697724448893373
Validation accuracy with LogisticRegression: 0.83
Regression on training set: 0.8466549736908043
Regression Score on test set: 0.8338762214983714
mae for LogisticRegression: 0.16612377850162866
mse for LogisticRegression: 0.16612377850162866
accuracy for LogisticRegression:0.8338762214983714
variance for LogisticRegression :0.014201230193749192

```

```

[231]:      features  estimatedCoefficients
1    Critic_Score          0.008154
0  Year_of_Release          0.002191

```

Confusion Matrix, Classification Report,ROC curve

```

[204]: from sklearn import metrics
from sklearn.metrics import roc_curve
from sklearn.metrics import auc

```

```

ypred = log_reg.predict(Xtest)
print("Confusion Matrix")
print(confusion_matrix(ytest, ypred))
cm = confusion_matrix(ytest, ypred)
np.set_printoptions(precision=2)
plt.figure()
plot_confusion_matrix(cm, classes=["Positive", "Negative"],
                      title='Confusion matrix')

print("Classification Report")
print(classification_report(ytest, ypred))
plot_classification_report(ytest, ypred)

logit_roc_auc = roc_auc_score(ytest, ypred)
print("Logistic AUC = %.2f" %logit_roc_auc)

b = log_reg.predict_proba(Xtest)[: ,1]
print(b[0:5])
fpr, tpr, threshold = roc_curve(ytest, b)

# plotting ROC curve
import matplotlib.pyplot as plt
plt.figure()
plt.plot(fpr, tpr, label='ROC curve (area = %.2f)' %logit_roc_auc)
plt.plot([0,1], [0,1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

```

Confusion Matrix

```
[[7257  37]
 [ 983  81]]
```

Confusion matrix, without normalization

```
[[7257  37]
 [ 983  81]]
```


Classification Report

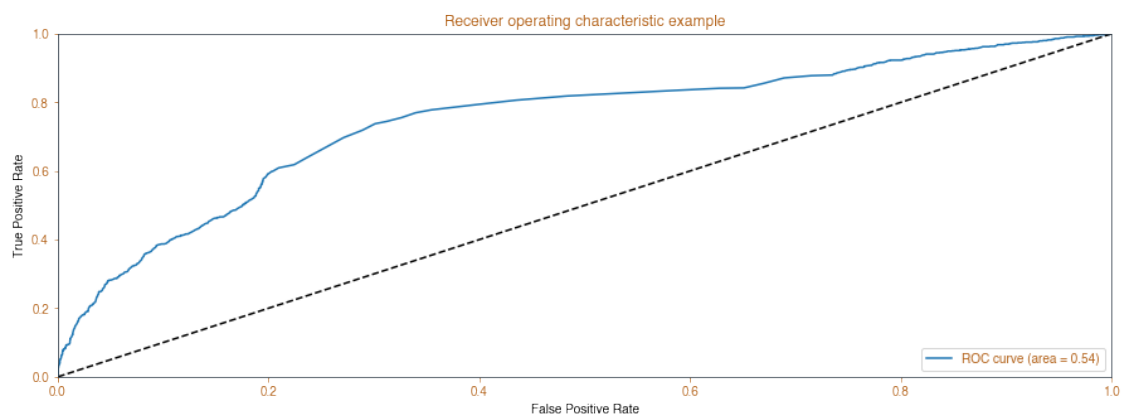

	precision	recall	f1-score	support
0	0.88	0.99	0.93	7294
1	0.69	0.08	0.14	1064
accuracy			0.88	8358
macro avg	0.78	0.54	0.54	8358

weighted avg 0.86 0.88 0.83 8358

Logistic AUC = 0.54

[0.1 0.09 0.05 0.16 0.1]






Data Analysis for Ratings. Encoding Ratings and clean up.

```
[205]: df = pd.read_csv('Video_Games_Sales.csv', encoding="utf-8")
df = df[df.Rating != "RP"]
df = df[df.Rating != "EC"]
df = df[df.Rating != "AO"]
df = df[df.Rating != "K-A"]

ax = df.groupby('Rating').sum().unstack().Global_Sales.
→ sort_values(ascending=False).head(10).plot(kind='bar',

→                                     figsize=(

→                                     13, 5));
```



```
[206]: df = pd.read_csv('Video_Games_Sales.csv', encoding="utf-8")
df = df[df.Platform != "2600"]
df = df[df.Platform != "3DO"]
df = df[df.Platform != "DC"]
df = df[df.Platform != "GEN"]
df = df[df.Platform != "GG"]
df = df[df.Platform != "NG"]
df = df[df.Platform != "PCFX"]
df = df[df.Platform != "SAT"]
df = df[df.Platform != "SCD"]
df = df[df.Platform != "TG16"]
df = df[df.Platform != "WS"]

df = df[df.User_Score != "tbd"]
df = df.dropna().reset_index(drop=True)
```

```

df["Platform_XB"] = 0
df["Platform_X360"] = 0
df["Platform_XOne"] = 0
df["Platform_PC"] = 0
df["Platform_PS"] = 0
df["Platform_PS2"] = 0
df["Platform_PS3"] = 0
df["Platform_PS4"] = 0
df["Platform_PSP"] = 0
df["Platform_PSV"] = 0
df["Platform_GB"] = 0
df["Platform_GBA"] = 0
df["Platform_DS"] = 0
df["Platform_3DS"] = 0
df["Platform_NES"] = 0
df["Platform_SNES"] = 0
df["Platform_N64"] = 0
df["Platform_GC"] = 0
df["Platform_Wii"] = 0
df["Platform_WiiU"] = 0

for elem in df.index.get_values():
    if df.get_value(elem, "Platform") == "XB": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "X360": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "XOne": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "PS": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "PS2": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "PS3": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "PS4": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "PSP": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "PSV": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "GB": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "GBA": df.set_value(elem, 1)
    if df.get_value(elem, "Platform") == "DS": df.set_value(elem, 1)

```

```

        if df.get_value(elem, "Platform") == "3DS": df.set_value(elem,
→"Platform_3DS", 1)
        if df.get_value(elem, "Platform") == "NES": df.set_value(elem,
→"Platform_NES", 1)
        if df.get_value(elem, "Platform") == "SNES": df.set_value(elem,
→"Platform_SNES", 1)
        if df.get_value(elem, "Platform") == "N64": df.set_value(elem,
→"Platform_N64", 1)
        if df.get_value(elem, "Platform") == "GC": df.set_value(elem,
→"Platform_GC", 1)
        if df.get_value(elem, "Platform") == "Wii": df.set_value(elem,
→"Platform_Wii", 1)
        if df.get_value(elem, "Platform") == "WiiU": df.set_value(elem,
→"Platform_WiiU", 1)

# Discretizzazione feature Genre

df["Genre_Action"] = 0
df["Genre_Adventure"] = 0
df["Genre_Fighting"] = 0
df["Genre_Misc"] = 0
df["Genre_Platform"] = 0
df["Genre_Puzzle"] = 0
df["Genre_Shooter"] = 0
df["Genre_Sports"] = 0
df["Genre_Simulation"] = 0
df["Genre_Strategy"] = 0
df["Genre_Racing"] = 0
df["Genre_Role-Playing"] = 0

for elem in df.index.get_values():
    if df.get_value(elem, "Genre") == "Action": df.set_value(elem,
→"Genre_Action", 1)
    if df.get_value(elem, "Genre") == "Adventure": df.set_value(elem,
→"Genre_Adventure", 1)
    if df.get_value(elem, "Genre") == "Fighting": df.set_value(elem,
→"Genre_Fighting", 1)
    if df.get_value(elem, "Genre") == "Misc": df.set_value(elem,
→"Genre_Misc", 1)
    if df.get_value(elem, "Genre") == "Platform": df.set_value(elem,
→"Genre_Platform", 1)
    if df.get_value(elem, "Genre") == "Puzzle": df.set_value(elem,
→"Genre_Puzzle", 1)
    if df.get_value(elem, "Genre") == "Shooter": df.set_value(elem,
→"Genre_Shooter", 1)

```

```

        if df.get_value(elem, "Genre") == "Sports": df.set_value(elem,
→"Genre_Sports", 1)
        if df.get_value(elem, "Genre") == "Simulation": df.set_value(elem,
→"Genre_Simulation", 1)
        if df.get_value(elem, "Genre") == "Strategy": df.set_value(elem,
→"Genre_Strategy", 1)
        if df.get_value(elem, "Genre") == "Racing": df.set_value(elem,
→"Genre_Racing", 1)
        if df.get_value(elem, "Genre") == "Role-Playing": df.set_value(elem,
→"Genre_Role-Playing", 1)

df["Rating_Everyone"] = 0
df["Rating_Everyone10"] = 0
df["Rating_Teen"] = 0
df["Rating_Mature"] = 0
df["Rating_Adult"] = 0

for elem in df.index.get_values():
    if df.get_value(elem, "Rating") == "E": df.set_value(elem,
→"Rating_Everyone", 1)
    if df.get_value(elem, "Rating") == "E10+": df.set_value(elem,
→"Rating_Everyone10", 1)
    if df.get_value(elem, "Rating") == "T": df.set_value(elem,
→"Rating_Teen", 1)
    if df.get_value(elem, "Rating") == "M": df.set_value(elem,
→"Rating_Mature", 1)
    if df.get_value(elem, "Rating") == "AO": df.set_value(elem,
→"Rating_Adult", 1)

# Discretizzazione feature Publisher

publisher_list = []

for elem in df.Publisher:
    if elem not in publisher_list:
        publisher_list.append(elem)

for elem in publisher_list:
    df[elem] = 0

for elem in df.index.get_values():
    df.set_value(elem, df.get_value(elem, "Publisher"), 1)

# Discretizzazione feature Developer

```

```

developer_list = []

for elem in df.Developer:
    if elem not in developer_list:
        developer_list.append(elem)

for elem in developer_list:
    df[elem] = 0

for elem in df.index.get_values():
    df.set_value(elem, df.get_value(elem, "Developer"), 1)

df = df[df.Rating != 'AO']
df = df[df.Rating != 'K-A']
df = df[df.Rating != 'RP']
df = df[df.Rating != 'EC']

del df['User_Score']
del df['User_Count']
del df['Critic_Score']
del df['Critic_Count']

del df['Platform']
del df['Genre']
del df['Publisher']
del df['Developer']

del df['Rating_Adult']

df = df.reindex(np.random.permutation(df.index)).reset_index(drop=True)

y_true = np.array(df['Rating'])
vector_names = np.array(df['Name'])
del df['Name']

y_true_int = np.empty(len(y_true), dtype=int)
i = 0
for elem in y_true:
    if elem == "E":
        y_true_int[i] = 1
    elif elem == "E10+":
        y_true_int[i] = 2
    elif elem == "T":
        y_true_int[i] = 3

```

```

        else:
            y_true_int[i] = 4
            i += 1

del df['Rating']

df = df.astype('float64')

print(" Dataset Analysis ")

print_e = df["Rating_Everyone"].value_counts()[1]
print_e10 = df["Rating_Everyone10"].value_counts()[1]
print_teen = df["Rating_Teen"].value_counts()[1]
print_mature = df["Rating_Mature"].value_counts()[1]

print("Number of elements: " + str(print_e + print_e10 + print_teen +
→print_mature) + "\n")

df_stampa = pd.DataFrame({"Rating": ['Everyone', 'Everyone 10+', 'Teen',
→'Mature'],
                           "Counts": [print_e, print_e10, print_teen,
→print_mature]})

cols = df_stampa.columns.tolist()
cols = cols[-1:] + cols[:-1]
df_stampa = df_stampa[cols]

print(df_stampa)

```

Dataset Analysis
Number of elements: 6808

	Counts	Rating
0	2079	Everyone
1	930	Everyone 10+
2	2367	Teen
3	1432	Mature

Utility Function for plot_learning_curve.

```

[207]: def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None, n_jobs=1,
→train_sizes=np.linspace(.1, 1.0, 5)):
    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")

```

```

plt.ylabel("Score")
train_sizes, train_scores, test_scores = learning_curve(estimator, X, y,
→cv=cv, n_jobs=n_jobs,

→train_sizes=train_sizes)
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)
plt.grid()

plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1, color="g")
plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="train_scores_mean")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Test_scores_mean")

plt.legend(loc="best")
return plt

```

[208]: *### 3. Third ML Problem*
X=> Year_of_Release Sales Platform
y=> RATING EVERYONE

[209]: *#### Random Forest*

[210]:

```

from sklearn.model_selection import learning_curve
import pandas as pd
import numpy as np
import seaborn as sns

from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import VotingClassifier

from sklearn.model_selection import learning_curve
import matplotlib.pyplot as plt

print("RATING EVERYONE")

```

```

df2 = df
y = df2['Rating_Everyone'].values
df2 = df2.drop(['Rating_Everyone'], axis=1)
df2 = df2.drop(['Rating_Everyone10'], axis=1)
df2 = df2.drop(['Rating_Teen'], axis=1)
df2 = df2.drop(['Rating_Mature'], axis=1)
X = df2.values

Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, test_size=0.20)

print("Random Forest: ")
plot_learning_curve(RandomForestClassifier(), "Rating_Everyone Learning_
→Curves (Random Forest)", X, y,
                    n_jobs=-1)

```

RATING EVERYONE
Random Forest:

[210]: <module 'matplotlib.pyplot' from
'/Users/nselvarajan/anaconda3/lib/python3.7/site-packages/matplotlib/pyplot.py'>




k-NN

[211]: `print("k-NN: ")`
`plot_learning_curve(KNeighborsClassifier(), "Rating_Everyone Learning Curves_`
`→(k-NN)", X, y, n_jobs=-1)`

k-NN:

[211]: <module 'matplotlib.pyplot' from
'/Users/nselvarajan/anaconda3/lib/python3.7/site-packages/matplotlib/pyplot.py'>



LogisticRegression

```
[232]: log_reg1 = LogisticRegression(penalty='l1', dual=False, C=1.0,
→fit_intercept=True, intercept_scaling=1,
                                class_weight=None, random_state=None,
→solver='liblinear', max_iter=100,
                                multi_class='ovr', verbose=0,
→warm_start=False, n_jobs=1)
    log_reg1.fit(Xtrain, ytrain)
    y_val_1 = log_reg1.predict(Xtest)
    ris = accuracy_score(ytest, y_val_1)
    mis = accuracy_score(ytest, y_val_1, normalize=False)

    print("Logistic Regression Rating_Everyone accuracy: ", ris)
    print("Logistic Regression Rating_Everyone misclassification: ", ytest.size
→- mis)
    MAE_rf = mean_absolute_error(ytest, y_val_1)
    MSE_rf = mean_squared_error(ytest, y_val_1)
    print("mae for Logistic Regression: " + str(MAE_rf))
    print("mse for Logistic Regression: " + str(MSE_rf))
    print("accuracy for Logistic Regression:" + str(log_reg1.
→score(Xtest, y_val_1)))
    print("variance for Logistic Regression : " +
→str(explained_variance_score(ytest, y_val_1)))
```

```
Logistic Regression Rating_Everyone accuracy:  0.8333750939614132
Logistic Regression Rating_Everyone misclassification:  665
mae for Logistic Regression: 0.16662490603858682
mse for Logistic Regression: 0.16662490603858682
```

accuracy for Logistic Regression:0.8333750939614132
variance for Logistic Regression :0.011810211222773925

RandomForestClassifier

```
[233]: radm1 = RandomForestClassifier(n_estimators=240, criterion='gini',  
    ↳max_depth=None, min_samples_split=2,  
        min_samples_leaf=1,  
        min_weight_fraction_leaf=0.0,  
    ↳max_features='auto', max_leaf_nodes=None,  
        min_impurity_split=1e-07, bootstrap=True,  
        oob_score=True, n_jobs=1, random_state=None,  
    ↳verbose=0, warm_start=False,  
        class_weight=None)  
  
    radm1.fit(Xtrain, ytrain)  
    y_val_l = radm1.predict(Xtest)  
    ris = accuracy_score(ytest, y_val_l)  
    mis = accuracy_score(ytest, y_val_l, normalize=False)  
    print("Random Forest Rating_Everyone accuracy: ", ris)  
    print("Random Forest Rating_Everyone misclassification: ", ytest.size - mis)  
    MAE_rf = mean_absolute_error(ytest, y_val_l)  
    MSE_rf = mean_squared_error(ytest, y_val_l)  
    print("mae for RandomForestClassifier: " + str(MAE_rf))  
    print("mse for RandomForestClassifier: " + str(MSE_rf))  
    print("accuracy for RandomForestClassifier:" + str(radm1.  
    ↳score(Xtest,ytest)))  
    print("variance for RandomForestClassifier : " +  
    ↳str(explained_variance_score(ytest, y_val_l)))
```

Random Forest Rating_Everyone accuracy: 0.8153345026309196
Random Forest Rating_Everyone misclassification: 737
mae for RandomForestClassifier: 0.18466549736908044
mse for RandomForestClassifier: 0.18466549736908044
accuracy for RandomForestClassifier:0.8153345026309196
variance for RandomForestClassifier :-0.24269676733883894

KNeighborsClassifier

```
[237]: knn1 = KNeighborsClassifier(n_neighbors=10, weights='uniform',  
    ↳algorithm='auto', leaf_size=30, p=1,  
        metric='minkowski', metric_params=None,  
    ↳n_jobs=1)  
  
    knn1.fit(Xtrain, ytrain)  
    y_val_l = knn1.predict(Xtest)  
    ris = accuracy_score(ytest, y_val_l)  
    mis = accuracy_score(ytest, y_val_l, normalize=False)  
    print("K-Nearest Neighbors Rating_Everyone accuracy: ", ris)
```

```

    print("K-Nearest Neighbors Rating_Everyone misclassification: ", ytest.size_
→- mis)
    print("K-Nearest Neighbors Rating_Everyone accuracy: ", ris)
    print("K-Nearest Neighbors Rating_Everyone misclassification: ", ytest.size_
→- mis)
    MAE_rf = mean_absolute_error(ytest, y_val_1)
    MSE_rf = mean_squared_error(ytest, y_val_1)
    print("mae for K-Nearest Neighbors: " + str(MAE_rf))
    print("mse for K-Nearest Neighbors: " + str(MSE_rf))
    print("accuracy for K-Nearest Neighbors:" + str(ris))
    print("variance for K-Nearest Neighbors :"+
→str(explained_variance_score(ytest, y_val_1)))

    print("\n")

```

```

K-Nearest Neighbors Rating_Everyone accuracy:  0.8376346780255575
K-Nearest Neighbors Rating_Everyone misclassification:  648
K-Nearest Neighbors Rating_Everyone accuracy:  0.8376346780255575
K-Nearest Neighbors Rating_Everyone misclassification:  648
mae for K-Nearest Neighbors: 0.1623653219744425
mse for K-Nearest Neighbors: 0.1623653219744425
accuracy for K-Nearest Neighbors:0.8376346780255575
variance for K-Nearest Neighbors :-0.03927400129078129

```

Evaluate Different Models

```

[215]: for clf, label in zip([log_reg1, radm1, knn1], ['Logistic Regression',
→'Random Forest', 'k-NN']):
        scores = cross_val_score(clf, X, y, cv=5, scoring='accuracy')
        print("Accuracy Score: %0.5f [%s]" % (scores.mean(), label))

```

```

Accuracy Score: 0.85165 [Logistic Regression]
Accuracy Score: 0.85899 [Random Forest]
Accuracy Score: 0.83681 [k-NN]

```

2.2 Comparison of Accuracy of the models

```

[216]: # # Comparison of Accuracy of the models
MSEs = [0.84768,0.85913,0.83519]

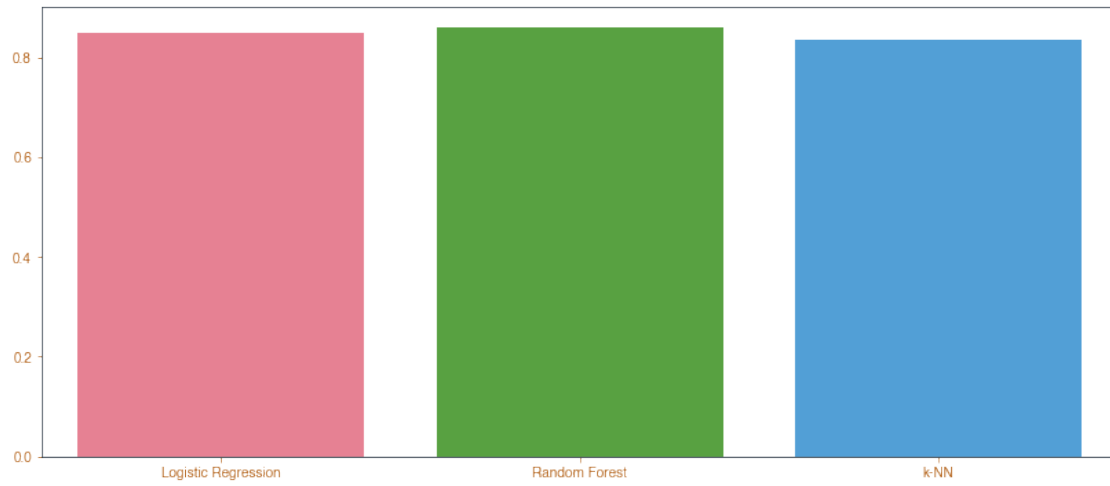
models = ['Logistic Regression','Random Forest','k-NN']
plt.figure(figsize=(14,6))
sns.barplot(models,MSEs,palette='husl')

```

```

[216]: <matplotlib.axes._subplots.AxesSubplot at 0x1a208b35f8>

```



[]:

[]:

[]: